

```

# Load necessary libraries
import pandas as pd
import numpy as np

# Load the dataset
data = pd.read_excel('/content/marketing_campaign_final.xlsx')

# Explore and preprocess data (handle missing values, outliers, encoding, etc.)
# Example preprocessing steps:
# data.info() # Check data types and missing values
# data.describe() # Summary statistics
# Handle missing values: data.fillna(), data.dropna()
# Encode categorical variables: pd.get_dummies(), LabelEncoder, etc.

from sklearn.preprocessing import LabelEncoder

# Assuming 'df' is your DataFrame and 'column_name' is the column you want to encode

# Create a LabelEncoder instance
label_encoder = LabelEncoder()

# Fit and transform the column to convert labels to numeric values
data['Education'] = label_encoder.fit_transform(data['Education'])

# Create a LabelEncoder instance
label_encoder = LabelEncoder()

# Fit and transform the column to convert labels to numeric values
data['Marital_Status'] = label_encoder.fit_transform(data['Marital_Status'])

data.drop(columns=['Dt_Customer'], axis=1, inplace=True)

data.fillna(0, inplace=True)

# Split the data into features (X) and target variable (y)
X = data.drop('Teenhome', axis=1) # Features
y = data['Teenhome'] # Target variable
print(data.head())

```

```

ID  Year_Birth  Education  Marital_Status  Income  Kidhome  Teenhome  \
0  5524      1957         2              4  58138.0      0         0
1  2174      1954         2              4  46344.0      1         1
2  4141      1965         2              5  71613.0      0         0
3  6182      1984         2              5  26646.0      1         0
4  5324      1981         4              3  58293.0      1         0

Recency  MntWines  MntFruits  ...  NumWebVisitsMonth  AcceptedCmp3  \
0      58      635      88  ...              7         0
1      38       11       1  ...              5         0
2      26      426      49  ...              4         0
3      26       11       4  ...              6         0
4      94      173      43  ...              5         0

AcceptedCmp4  AcceptedCmp5  AcceptedCmp1  AcceptedCmp2  Complain  \
0         0         0         0         0         0
1         0         0         0         0         0
2         0         0         0         0         0
3         0         0         0         0         0
4         0         0         0         0         0

Z_CostContact  Z_Revenue  Response
0             3         11         1
1             3         11         0
2             3         11         0

```

| | | | |
|---|---|----|---|
| 3 | 3 | 11 | 0 |
| 4 | 3 | 11 | 0 |

[5 rows x 28 columns]


```
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# Load the dataset
data = pd.read_excel('/content/marketing_campaign_final.xlsx')

# Example visualizations


# Histogram
sns.histplot(data['MntWines'], bins=20, kde=True)
plt.title('Histogram')

plt.show()
```




```
# Bar Plot
sns.countplot(x='Education', data=data)
plt.title('Bar Plot')
```

```
Text(0.5, 1.0, 'Bar Plot')
```



```
# Scatter Plot
sns.scatterplot(x='NumWebVisitsMonth', y='Income', data=data)
plt.title('Scatter Plot')
```

```
Text(0.5, 1.0, 'Scatter Plot')
```



```
# Pie Chart
data['Marital_Status'].value_counts().plot(kind='pie', autopct='%1.1f%%')
plt.title('Pie Chart')
```

```
Text(0.5, 1.0, 'Pie Chart')
```

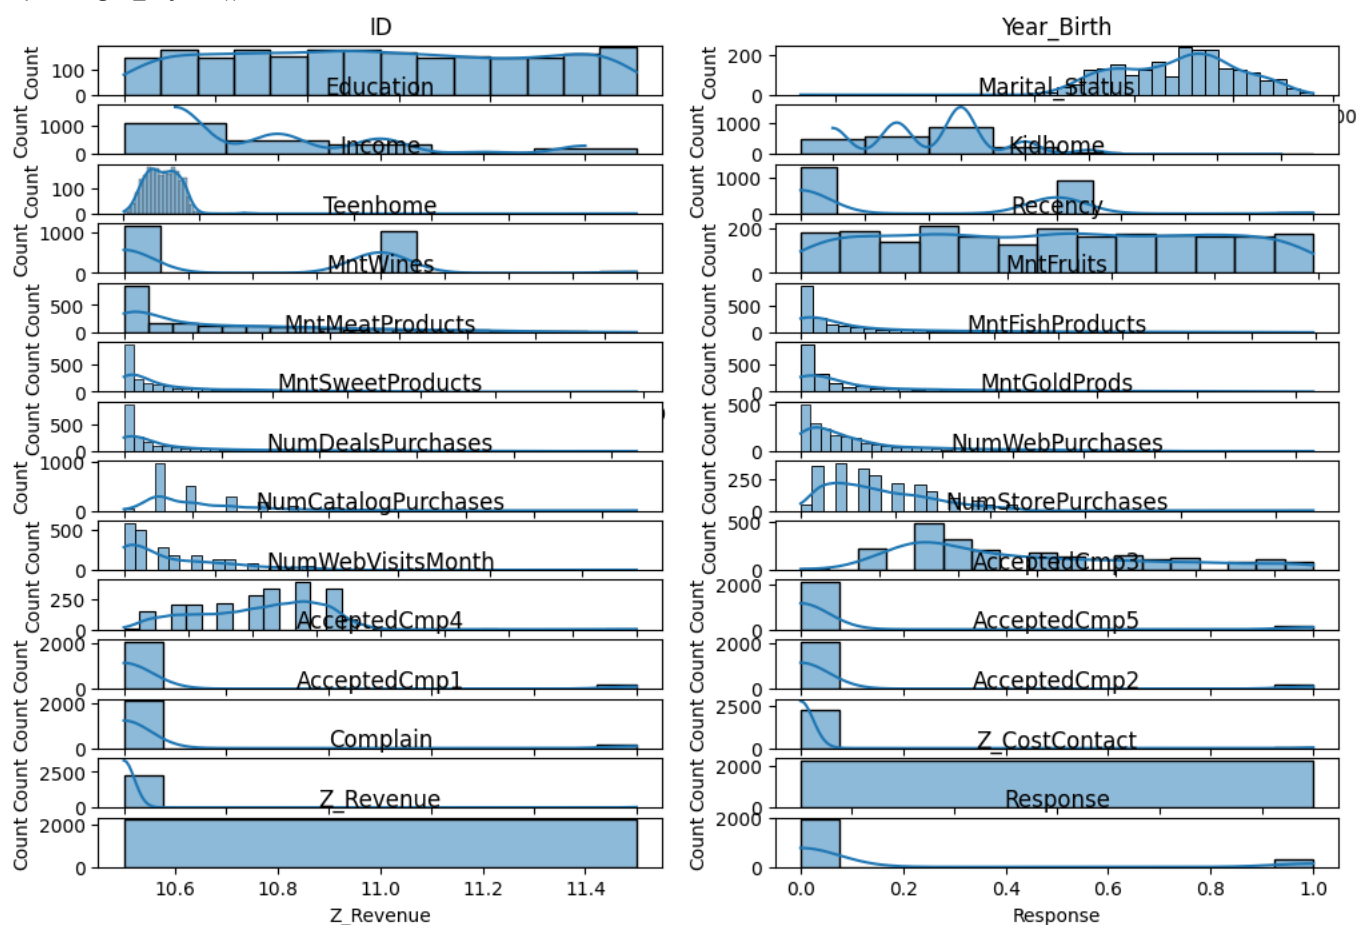
Pie Chart



```
# Histograms for each column in the dataset
data.drop('Dt_Customer', axis=1, inplace=True) # axis=1 indicates column-wise operation
```


```
plt.figure(figsize=(12, 8))
num_columns = len(data.columns)
for i, column in enumerate(data.columns):
    plt.subplot((num_columns // 2) + (num_columns % 2), 2, i + 1)
    sns.histplot(data[column], kde=True)
    plt.title(column)
plt.tight_layout()
plt.show()
```

```
<ipython-input-54-27df3b85dbc4>:10: UserWarning: Tight layout not applied. tight_layout cannot make axes height s
plt.tight_layout()
```



```
# Box Plot (Interquartile Range)
sns.boxplot(x='Kidhome', y='Recency', data=data)
plt.title('Box Plot')
```

```
Text(0.5, 1.0, 'Box Plot')
```



```
# Check the first few rows of the DataFrame
print(data.head())
```

```
# List all columns in the DataFrame
print(data.columns)
```

```
# Check the shape of the DataFrame
print(data.shape)
```

| | ID | Year_Birth | Education | Marital_Status | Income | Kidhome | Teenhome | \ |
|---|------|------------|------------|----------------|---------|---------|----------|---|
| 0 | 5524 | 1957 | Graduation | Single | 58138.0 | 0 | 0 | |
| 1 | 2174 | 1954 | Graduation | Single | 46344.0 | 1 | 1 | |
| 2 | 4141 | 1965 | Graduation | Together | 71613.0 | 0 | 0 | |
| 3 | 6182 | 1984 | Graduation | Together | 26646.0 | 1 | 0 | |
| 4 | 5324 | 1981 | PhD | Married | 58293.0 | 1 | 0 | |

| | Recency | MntWines | MntFruits | ... | NumWebVisitsMonth | AcceptedCmp3 | \ |
|---|---------|----------|-----------|-----|-------------------|--------------|---|
| 0 | 58 | 635 | 88 | ... | 7 | 0 | |
| 1 | 38 | 11 | 1 | ... | 5 | 0 | |
| 2 | 26 | 426 | 49 | ... | 4 | 0 | |
| 3 | 26 | 11 | 4 | ... | 6 | 0 | |
| 4 | 94 | 173 | 43 | ... | 5 | 0 | |

| | AcceptedCmp4 | AcceptedCmp5 | AcceptedCmp1 | AcceptedCmp2 | Complain | \ |
|---|--------------|--------------|--------------|--------------|----------|---|
| 0 | 0 | 0 | 0 | 0 | 0 | |
| 1 | 0 | 0 | 0 | 0 | 0 | |
| 2 | 0 | 0 | 0 | 0 | 0 | |
| 3 | 0 | 0 | 0 | 0 | 0 | |
| 4 | 0 | 0 | 0 | 0 | 0 | |

| | Z_CostContact | Z_Revenue | Response |
|---|---------------|-----------|----------|
| 0 | 3 | 11 | 1 |
| 1 | 3 | 11 | 0 |
| 2 | 3 | 11 | 0 |
| 3 | 3 | 11 | 0 |
| 4 | 3 | 11 | 0 |

```
[5 rows x 28 columns]
```

```
Index(['ID', 'Year_Birth', 'Education', 'Marital_Status', 'Income', 'Kidhome',
      'Teenhome', 'Recency', 'MntWines', 'MntFruits', 'MntMeatProducts',
      'MntFishProducts', 'MntSweetProducts', 'MntGoldProds',
      'NumDealsPurchases', 'NumWebPurchases', 'NumCatalogPurchases',
      'NumStorePurchases', 'NumWebVisitsMonth', 'AcceptedCmp3',
      'AcceptedCmp4', 'AcceptedCmp5', 'AcceptedCmp1', 'AcceptedCmp2',
      'Complain', 'Z_CostContact', 'Z_Revenue', 'Response'],
      dtype='object')
```

```
(2240, 28)
```

```
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegressionCV

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Using L1 regularization
logreg = LogisticRegression(max_iter=10000, solver='saga', penalty='l1' , tol=1e-3)
logreg.fit(X_train, y_train)

# Predict on the test set
predictions = logreg.predict(X_test)

# Evaluate the model
accuracy = logreg.score(X_test, y_test)
accuracy_percentage = accuracy * 100
print(f"Accuracy: {accuracy_percentage:.2f}%")

    Accuracy: 86.00%

from sklearn.linear_model import LogisticRegression
from sklearn.datasets import load_iris # Using Iris dataset as an example

# Load the dataset (replace this with your dataset)
data = load_iris()
X = data.data # Features
y = data.target # Target

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the Logistic Regression model
logreg = LogisticRegression(max_iter=1000 , solver='saga', penalty='l1' , tol=1e-3) # You can specify other paramete

# Fit the model on the training data
logreg.fit(X_train, y_train)

# Predict on the test set
predictions = logreg.predict(X_test)

# Evaluate the model
accuracy = logreg.score(X_test, y_test)
print(f"Accuracy: {accuracy:.2f}")

    Accuracy: 1.00
```

```

from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.datasets import load_iris # Using Iris dataset as an example

# Load the dataset (replace this with your dataset)
data = load_iris()
X = data.data # Features
y = data.target # Target

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the Support Vector Classifier
svm = SVC(kernel='rbf') # You can specify other parameters like kernel, C, gamma, etc.

# Fit the model on the training data
svm.fit(X_train, y_train)

# Predict on the test set
predictions = svm.predict(X_test)

# Evaluate the model
accuracy = svm.score(X_test, y_test)
accuracy_percentage = accuracy * 100
print(f"Accuracy: {accuracy_percentage:.2f}%")

    Accuracy: 100.00%

from sklearn.decomposition import PCA
from sklearn.svm import SVC

# Load the dataset (replace this with your dataset)
data = load_iris()
X = data.data[:, :2] # Considering only the first two features for visualization purposes
y = data.target

# Initialize the SVM classifier
svm = SVC(kernel='linear') # You can use different kernels

# Fit the SVM on the data
svm.fit(X, y)


# Use PCA to reduce dimensions for visualization (for plotting purposes)
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X)

# Create a meshgrid to plot decision boundary
x_min, x_max = X_pca[:, 0].min() - 1, X_pca[:, 0].max() + 1
y_min, y_max = X_pca[:, 1].min() - 1, X_pca[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.1),
                     np.arange(y_min, y_max, 0.1))

# Get predictions for meshgrid points
Z = svm.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

# Plot decision boundary
plt.contourf(xx, yy, Z, alpha=0.4)
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=y, edgecolor='k')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.title('SVM Decision Boundary Visualization (PCA)')
plt.show()

```



```
from sklearn.tree import DecisionTreeClassifier
# Using Iris dataset as an example

# Load the dataset (replace this with your dataset)
data = load_iris()
X = data.data # Features
y = data.target # Target

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the Decision Tree Classifier
clf = DecisionTreeClassifier(max_depth=3, random_state=42) # You can specify other parameters like max_depth, criterion

# Fit the model on the training data
clf.fit(X_train, y_train)

# Predict on the test set
predictions = clf.predict(X_test)

# Evaluate the model
accuracy = clf.score(X_test, y_test)
accuracy_percentage = accuracy * 100
print(f"Accuracy: {accuracy_percentage:.2f}")

Accuracy: 100.00
```



```

from sklearn.tree import DecisionTreeClassifier, plot_tree
import matplotlib.pyplot as plt

# Load the dataset (replace this with your dataset)
data = load_iris()
X = data.data # Features
y = data.target # Target

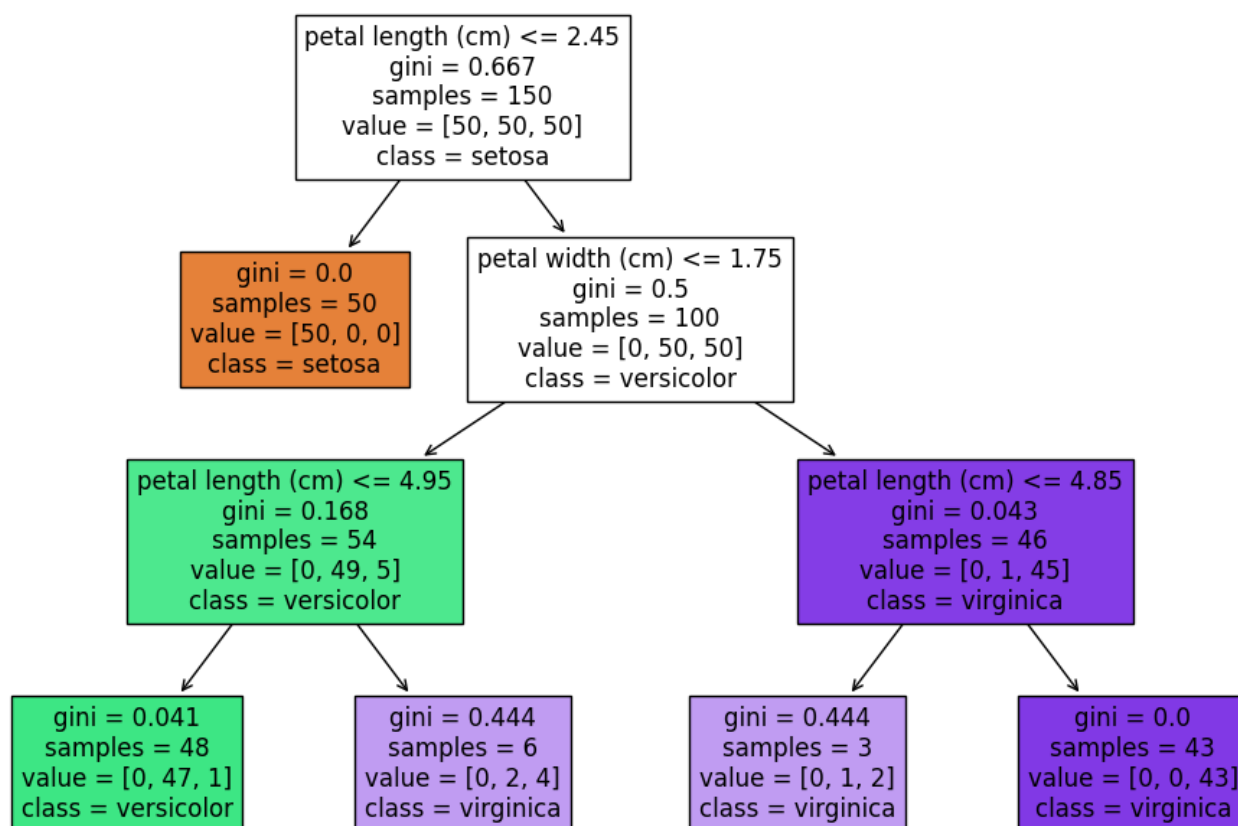
# Initialize the Decision Tree Classifier
clf = DecisionTreeClassifier(max_depth=3, random_state=42) # You can specify other parameters like max_depth, criter

# Fit the model on the entire dataset (for demonstration purposes)
clf.fit(X, y)

# Plot the Decision Tree
plt.figure(figsize=(12, 8))
plot_tree(clf, filled=True, feature_names=data.feature_names, class_names=data.target_names)
plt.title("Decision Tree Visualization")
plt.show()

```

Decision Tree Visualization



```

from sklearn.datasets import fetch_california_housing
from sklearn.tree import DecisionTreeRegressor, plot_tree
import matplotlib.pyplot as plt

# Load the California housing dataset
data = fetch_california_housing()
X = data.data # Features
y = data.target # Target (continuous variable for regression)

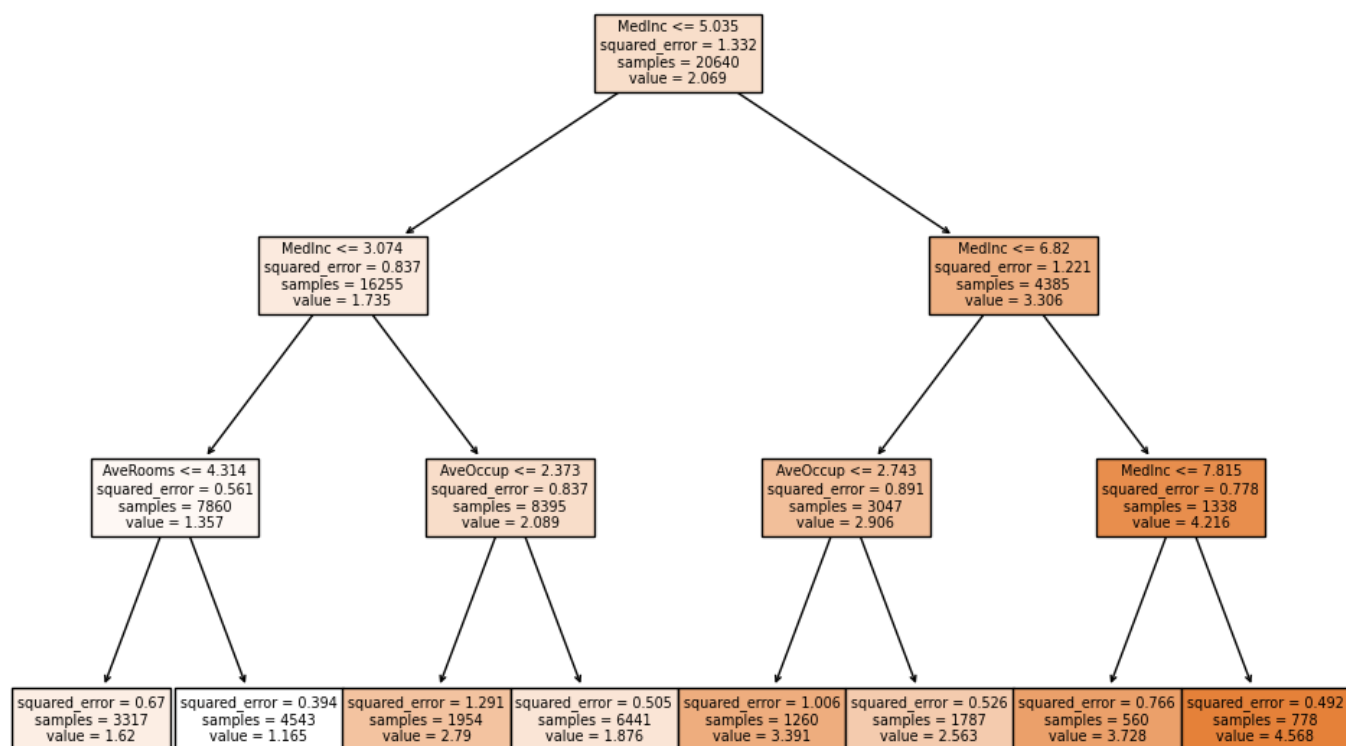
# Initialize the CART Regressor
regressor = DecisionTreeRegressor(max_depth=3, random_state=42) # You can specify other parameters like max_depth, c

# Fit the model on the entire dataset (for demonstration purposes)
regressor.fit(X, y)

# Plot the CART Decision Tree for regression
plt.figure(figsize=(12, 8))
plot_tree(regressor, filled=True, feature_names=data.feature_names)
plt.title("CART Decision Tree Visualization (Regression)")
plt.show()

```

CART Decision Tree Visualization (Regression)



```
from sklearn.datasets import make_classification
from sklearn.metrics import roc_curve, roc_auc_score

# Generate synthetic data for demonstration purposes
X, y = make_classification(n_samples=1000, n_features=20, n_classes=2, random_state=42)

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Initialize the Logistic Regression model
model = LogisticRegression()

# Fit the model on the training data
model.fit(X_train, y_train)

# Predict probabilities for the test set
y_prob = model.predict_proba(X_test)[:, 1]

# Calculate ROC curve and AUC score
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
auc = roc_auc_score(y_test, y_prob)

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f'AUC = {auc:.2f}')
plt.plot([0, 1], [0, 1], 'k--') # Diagonal line representing random guessing
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()
plt.show()
```



```
# Generate synthetic data for demonstration purposes
X, y = make_classification(n_samples=1000, n_features=20, n_classes=2, random_state=42)

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the Random Forest Classifier
rf = RandomForestClassifier(n_estimators=100, random_state=42) # You can specify other parameters like max_depth, et

# Fit the model on the training data
rf.fit(X_train, y_train)

# Predict on the test set
predictions = rf.predict(X_test)

# Evaluate accuracy
accuracy = accuracy_score(y_test, predictions)
accuracy_percentage = accuracy*100
print(f"Random Forest Accuracy: {accuracy_percentage:.2f}%")

Random Forest Accuracy: 90.00%
```

```
from sklearn.datasets import make_classification
from sklearn.ensemble import GradientBoostingClassifier

# Generate synthetic data for demonstration purposes
X, y = make_classification(n_samples=1000, n_features=20, n_classes=2, random_state=42)

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the Gradient Boosting Classifier
gb = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1, random_state=42) # You can specify other parame

# Fit the model on the training data
gb.fit(X_train, y_train)

# Predict on the test set
predictions = gb.predict(X_test)

# Evaluate accuracy
accuracy = accuracy_score(y_test, predictions)
accuracy_percentage = accuracy*100;
print(f"Gradient Boosting Accuracy: {accuracy_percentage:.2f}%")

Gradient Boosting Accuracy: 91.00%
```

```
# Generate synthetic data for demonstration purposes
X, y = make_classification(n_samples=1000, n_features=20, n_classes=2, random_state=42)

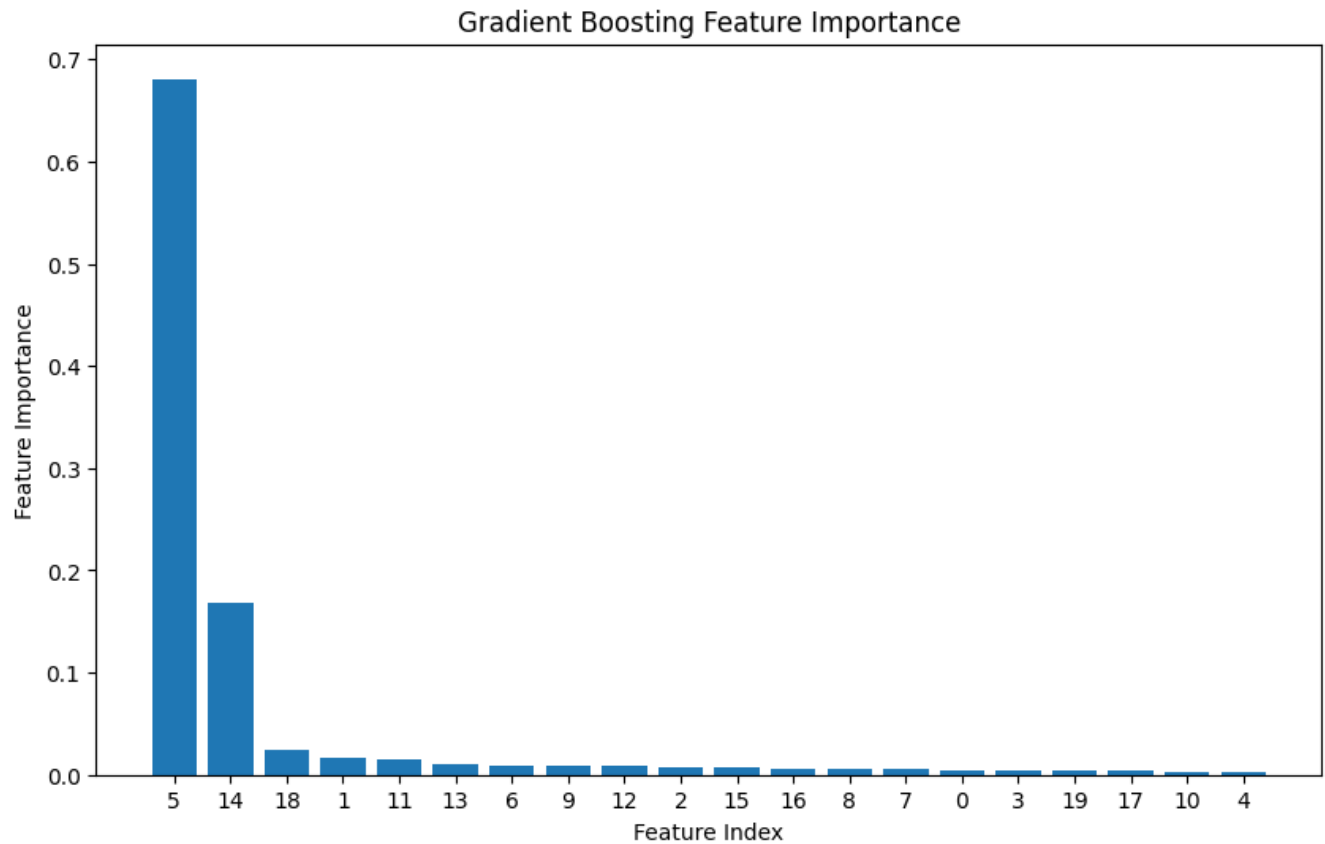
# Initialize the Gradient Boosting Classifier
gb = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1, random_state=42)
gb.fit(X, y)

# Get feature importances
feature_importance = gb.feature_importances_

# Sort feature importances in descending order
indices = feature_importance.argsort()[::-1]

# Plot feature importances
plt.figure(figsize=(10, 6))
plt.title("Gradient Boosting Feature Importance")
```

```
plt.bar(range(X.shape[1]), feature_importance[indices], align="center")
plt.xticks(range(X.shape[1]), indices)
plt.xlabel("Feature Index")
plt.ylabel("Feature Importance")
plt.show()
```



```
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.datasets import load_iris # Using Iris dataset as an example

# Load the dataset (replace this with your dataset)
data = load_iris()
X = data.data # Features
y = data.target # Target

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the Random Forest Classifier
rf = RandomForestClassifier(n_estimators=100, random_state=42) # You can specify other parameters like n_estimators,

# Fit the model on the training data
rf.fit(X_train, y_train)

# Predict on the test set
predictions = rf.predict(X_test)

# Evaluate the model
accuracy = rf.score(X_test, y_test)
accuracy_percentage = accuracy * 100
print(f"Accuracy: {accuracy_percentage:.2f}%")
```

Accuracy: 100.00%

```
# Create NumPy arrays from the columns
recency = np.array(data['Recency'])
kidhome = np.array(data['Kidhome'])

# Compute Pearson correlation coefficient using NumPy
correlation_matrix = np.corrcoef(recency, kidhome)

# Extract the correlation coefficient
correlation = correlation_matrix[0, 1]

print(f"Pearson Correlation for recency and kidhome: {correlation}")

    Pearson Correlation for recency and kidhome: -1.2819751242557092e-17

print(data.columns)

    Index(['ID', 'Year_Birth', 'Education', 'Marital_Status', 'Income', 'Kidhome',
    'Teenhome', 'Recency', 'MntWines', 'MntFruits', 'MntMeatProducts',
    'MntFishProducts', 'MntSweetProducts', 'MntGoldProds',
    'NumDealsPurchases', 'NumWebPurchases', 'NumCatalogPurchases',
    'NumStorePurchases', 'NumWebVisitsMonth', 'AcceptedCmp3',
    'AcceptedCmp4', 'AcceptedCmp5', 'AcceptedCmp1', 'AcceptedCmp2',
    'Complain', 'Z_CostContact', 'Z_Revenue', 'Response'],
    dtype='object')
```

```
# Basic filter method: Filtering rows based on a condition
# For example, filtering rows where 'Column_A' values are greater than 50
filtered_df = data[data['Recency'] > 98]

# Display the filtered DataFrame
print("\nFiltered DataFrame:")
print(filtered_df)
```

Filtered DataFrame:

| | ID | Year_Birth | Education | Marital_Status | Income | Kidhome | \ |
|------|-------|------------|------------|----------------|---------|---------|---|
| 38 | 8595 | 1973 | Graduation | Widow | 42429.0 | 0 | |
| 192 | 7829 | 1900 | 2n Cycle | Divorced | 36640.0 | 1 | |
| 208 | 868 | 1966 | Graduation | Married | 44794.0 | 0 | |
| 444 | 2106 | 1974 | 2n Cycle | Married | 20130.0 | 0 | |
| 491 | 22 | 1976 | Graduation | Divorced | 46310.0 | 1 | |
| 606 | 7232 | 1973 | Graduation | Widow | 42429.0 | 0 | |
| 685 | 10142 | 1976 | PhD | Divorced | 66476.0 | 0 | |
| 700 | 9977 | 1973 | Graduation | Divorced | 78901.0 | 0 | |
| 725 | 7212 | 1966 | Graduation | Married | 44794.0 | 0 | |
| 1033 | 5263 | 1977 | 2n Cycle | Married | 31056.0 | 1 | |
| 1171 | 3363 | 1974 | 2n Cycle | Married | 20130.0 | 0 | |
| 1473 | 4070 | 1969 | PhD | Married | 94871.0 | 0 | |
| 1542 | 528 | 1978 | Graduation | Married | 65819.0 | 0 | |
| 1685 | 7947 | 1969 | Graduation | Married | 42231.0 | 1 | |
| 1800 | 2831 | 1976 | Graduation | Together | 78416.0 | 0 | |
| 1820 | 2415 | 1962 | Graduation | Together | 62568.0 | 0 | |
| 1894 | 1743 | 1974 | Graduation | Single | 69719.0 | 0 | |

| | Teenhome | Recency | MntWines | MntFruits | ... | NumWebVisitsMonth | \ |
|------|----------|---------|----------|-----------|-----|-------------------|---|
| 38 | 1 | 99 | 55 | 0 | ... | 5 | |
| 192 | 0 | 99 | 15 | 6 | ... | 5 | |
| 208 | 1 | 99 | 54 | 0 | ... | 6 | |
| 444 | 0 | 99 | 0 | 6 | ... | 8 | |
| 491 | 0 | 99 | 185 | 2 | ... | 8 | |
| 606 | 1 | 99 | 55 | 0 | ... | 5 | |
| 685 | 1 | 99 | 372 | 18 | ... | 4 | |
| 700 | 1 | 99 | 321 | 11 | ... | 4 | |
| 725 | 1 | 99 | 54 | 0 | ... | 6 | |
| 1033 | 0 | 99 | 5 | 10 | ... | 8 | |
| 1171 | 0 | 99 | 0 | 6 | ... | 8 | |
| 1473 | 2 | 99 | 169 | 24 | ... | 7 | |

| | | | | | | |
|------|---|----|-----|----|-----|---|
| 1542 | 0 | 99 | 267 | 38 | ... | 3 |
| 1685 | 1 | 99 | 24 | 0 | ... | 5 |
| 1800 | 1 | 99 | 453 | 38 | ... | 3 |
| 1820 | 1 | 99 | 362 | 17 | ... | 4 |
| 1894 | 0 | 99 | 273 | 86 | ... | 1 |

| | AcceptedCmp3 | AcceptedCmp4 | AcceptedCmp5 | AcceptedCmp1 | AcceptedCmp2 | \ |
|------|--------------|--------------|--------------|--------------|--------------|---|
| 38 | 0 | 0 | 0 | 0 | 0 | 0 |
| 192 | 0 | 0 | 0 | 0 | 0 | 0 |
| 208 | 0 | 0 | 0 | 0 | 0 | 0 |
| 444 | 0 | 0 | 0 | 0 | 0 | 0 |
| 491 | 0 | 0 | 0 | 0 | 0 | 0 |
| 606 | 0 | 0 | 0 | 0 | 0 | 0 |
| 685 | 0 | 0 | 0 | 0 | 0 | 0 |
| 700 | 0 | 0 | 0 | 0 | 0 | 0 |
| 725 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1033 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1171 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1473 | 0 | 1 | 1 | 0 | 0 | 0 |
| 1542 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1685 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1800 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1820 | 0 | 0 | 0 | 1 | 0 | 0 |
| 1894 | 0 | 0 | 0 | 0 | 0 | 0 |

```

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Assuming 'X' contains your features and 'y' contains the target variable

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the Logistic Regression model
logreg = LogisticRegression(max_iter=1000) # You can adjust max_iter as needed

# Train the model on the training data
logreg.fit(X_train, y_train)

# Predict on the test set
y_pred = logreg.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")

# Print classification report and confusion matrix
print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

```

```

Accuracy: 1.00
Classification Report:
              precision    recall  f1-score   support

     0       1.00      1.00      1.00        10
     1       1.00      1.00      1.00         9
     2       1.00      1.00      1.00        11

   accuracy                   1.00        30
  macro avg       1.00      1.00      1.00        30
 weighted avg       1.00      1.00      1.00        30

```

```

Confusion Matrix:
[[10  0  0]

```

```
[ 0  9  0]  
[ 0  0 11]
```

```
df = pd.DataFrame(data)  
  
# Compute Pearson correlation matrix  
corr_matrix = df.corr()  
  
# Create a heatmap using seaborn  
plt.figure(figsize=(8, 6))  
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt='.2f')  
plt.title('Pearson Correlation Heatmap')  
plt.show()
```

Pearson Correlation Heatmap



Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.