

```

# Load necessary libraries
import pandas as pd
import numpy as np
import plotly.express as px
# Load the dataset
data = pd.read_excel('/content/marketing_campaign_final.xlsx')

# Explore and preprocess data (handle missing values, outliers, encoding, etc.)
# Example preprocessing steps:
# data.info() # Check data types and missing values
# data.describe() # Summary statistics
# Handle missing values: data.fillna(), data.dropna()
# Encode categorical variables: pd.get_dummies(), LabelEncoder, etc.

from sklearn.preprocessing import LabelEncoder

# Assuming 'df' is your DataFrame and 'column_name' is the column you want to encode

# Create a LabelEncoder instance
label_encoder = LabelEncoder()

# Fit and transform the column to convert labels to numeric values
data['Education'] = label_encoder.fit_transform(data['Education'])

# Create a LabelEncoder instance
label_encoder = LabelEncoder()

# Fit and transform the column to convert labels to numeric values
data['Marital_Status'] = label_encoder.fit_transform(data['Marital_Status'])

data.drop(columns=['Dt_Customer'], axis=1, inplace=True)

data.fillna(0, inplace=True)

# Split the data into features (X) and target variable (y)
X = data.drop('Teenhome', axis=1) # Features
y = data['Teenhome'] # Target variable
print(data.head())

```

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	Teenhome	\
0	5524	1957	2	4	58138.0	0	0	
1	2174	1954	2	4	46344.0	1	1	
2	4141	1965	2	5	71613.0	0	0	
3	6182	1984	2	5	26646.0	1	0	
4	5324	1981	4	3	58293.0	1	0	

	Recency	MntWines	MntFruits	...	NumWebVisitsMonth	AcceptedCmp3	\
0	58	635	88	...	7	0	
1	38	11	1	...	5	0	
2	26	426	49	...	4	0	
3	26	11	4	...	6	0	
4	94	173	43	...	5	0	

	AcceptedCmp4	AcceptedCmp5	AcceptedCmp1	AcceptedCmp2	Complain	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	0	

	Z_CostContact	Z_Revenue	Response
0	3	11	1
1	3	11	0
2	3	11	0

```

3          3          11          0
4          3          11          0

```

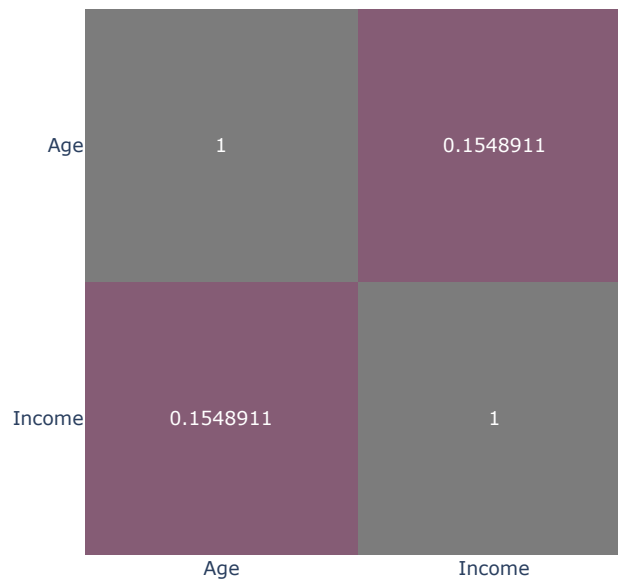
```
[5 rows x 28 columns]
```

```

data['Age'] = 2023 - data['Year_Birth']
px.imshow(data[["Age", "Income"]].corr(), text_auto=True,
           title="Correlation Between Age and Income", color_continuous_scale=px.colors.qualitative.Antique)

```

Correlation Between Age and Income



```

## Spending feature creation
data['spending'] = data['MntWines'] + data['MntFruits'] + data['MntMeatProducts'] + data['MntFishProducts'] \
                  + data['MntSweetProducts'] + data['MntGoldProds']

## rename some cols
data=data.rename(columns={'NumWebPurchases': 'Web', 'NumCatalogPurchases': 'Catalog', 'NumStorePurchases': 'Store'})


used_cols = ["Kidhome", "Teenhome"]

fig, ax = plt.subplots(1, 2, figsize = (18,5))

for i in range(2):
    sns.countplot(ax=ax[i], data=data, x=used_cols[i])
    ax[i].set_title(f"{used_cols[i]} Countplot.", weight="bold")
    ax[i].set_ylabel("Count")

plt.show()

```



```
## Children feature creation
data['Children'] = data['Kidhome'] + data['Teenhome']

def has_chid_or_no(x):
    if x > 0:
        return "Has child"
    else :
        return "no child"

data['Has_child'] = data["Children"].apply(has_chid_or_no)
data['Has_child'].value_counts()

    Has child    1602
    no child     638
    Name: Has_child, dtype: int64

## drop outliers
data = data.loc[(data["Income"] <= 200000)]
## reset index after drop outliers
data.reset_index(drop=True, inplace=True)
print(f'Shape After Drop Outliers : {data.shape}')
print(f'Statistics Info. about Income : \n{data["Income"].describe()}')

    Shape After Drop Outliers : (2239, 32)
    Statistics Info. about Income :
    count      2239.000000
    mean       51412.792765
    std        22069.582225
    min         0.000000
    25%        34716.000000
    50%        51039.000000
    75%        68277.500000
    max       162397.000000
    Name: Income, dtype: float64


import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# Load the dataset
data = pd.read_excel('/content/marketing_campaign_final.xlsx')

# Example visualizations


# Histogram
sns.histplot(data['MntWines'], bins=20, kde=True)
plt.title('Histogram')

plt.show()
```




```
# Bar Plot
sns.countplot(x='Education', data=data)
plt.title('Bar Plot')
```

```
Text(0.5, 1.0, 'Bar Plot')
```



```
# Scatter Plot
sns.scatterplot(x='NumWebVisitsMonth', y='Income', data=data)
plt.title('Scatter Plot')
```


```
Text(0.5, 1.0, 'Scatter Plot')
```



```
# Pie Chart
```

```
data['Marital_Status'].value_counts().plot(kind='pie', autopct='%1.1f%%')
plt.title('Pie Chart')
```

```
Text(0.5, 1.0, 'Pie Chart')
```

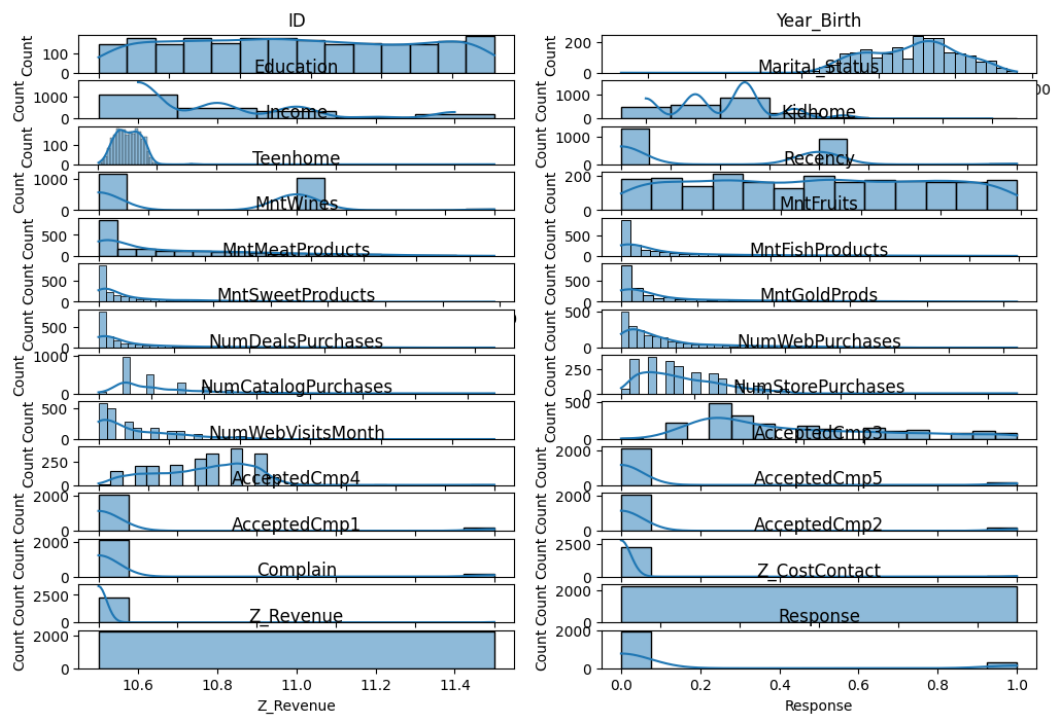


```
# Histograms for each column in the dataset
```

```
data.drop('Dt_Customer', axis=1, inplace=True) # axis=1 indicates column-wise operation
```


```
plt.figure(figsize=(12, 8))
num_columns = len(data.columns)
for i, column in enumerate(data.columns):
    plt.subplot((num_columns // 2) + (num_columns % 2), 2, i + 1)
    sns.histplot(data[column], kde=True)
    plt.title(column)
plt.tight_layout()
plt.show()
```

```
<ipython-input-54-27df3b85dbc4>:10: UserWarning: Tight layout not applied. tight_layout
plt.tight_layout()
```



```
# Box Plot (Interquartile Range)
sns.boxplot(x='Kidhome', y='Recency', data=data)
plt.title('Box Plot')
```

```
Text(0.5, 1.0, 'Box Plot')
```



```
# Check the first few rows of the DataFrame
print(data.head())
```

```
# List all columns in the DataFrame
print(data.columns)
```

```
# Check the shape of the DataFrame
print(data.shape)
```

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	Teenhome	\
0	5524	1957	Graduation	Single	58138.0	0	0	
1	2174	1954	Graduation	Single	46344.0	1	1	
2	4141	1965	Graduation	Together	71613.0	0	0	
3	6182	1984	Graduation	Together	26646.0	1	0	
4	5324	1981	PhD	Married	58293.0	1	0	

	Recency	MntWines	MntFruits	...	NumWebVisitsMonth	AcceptedCmp3	\
0	58	635	88	...	7	0	
1	38	11	1	...	5	0	
2	26	426	49	...	4	0	
3	26	11	4	...	6	0	
4	94	173	43	...	5	0	

	AcceptedCmp4	AcceptedCmp5	AcceptedCmp1	AcceptedCmp2	Complain	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	0	

	Z_CostContact	Z_Revenue	Response
0	3	11	1
1	3	11	0
2	3	11	0
3	3	11	0
4	3	11	0

```
[5 rows x 28 columns]
```

```
Index(['ID', 'Year_Birth', 'Education', 'Marital_Status', 'Income', 'Kidhome',
      'Teenhome', 'Recency', 'MntWines', 'MntFruits', 'MntMeatProducts',
      'MntFishProducts', 'MntSweetProducts', 'MntGoldProds',
      'NumDealsPurchases', 'NumWebPurchases', 'NumCatalogPurchases',
      'NumStorePurchases', 'NumWebVisitsMonth', 'AcceptedCmp3',
      'AcceptedCmp4', 'AcceptedCmp5', 'AcceptedCmp1', 'AcceptedCmp2',
      'Complain', 'Z_CostContact', 'Z_Revenue', 'Response'],
      dtype='object')
(2240, 28)
```

```
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegressionCV

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Using L1 regularization
logreg = LogisticRegression(max_iter=10000, solver='saga', penalty='l1' , tol=1e-3)
logreg.fit(X_train, y_train)

# Predict on the test set
predictions = logreg.predict(X_test)

# Evaluate the model
accuracy = logreg.score(X_test, y_test)
accuracy_percentage = accuracy * 100
print(f"Accuracy: {accuracy_percentage:.2f}%")

    Accuracy: 86.00%

from sklearn.linear_model import LogisticRegression
from sklearn.datasets import load_iris # Using Iris dataset as an example

# Load the dataset (replace this with your dataset)
data = load_iris()
X = data.data # Features
y = data.target # Target

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the Logistic Regression model
logreg = LogisticRegression(max_iter=1000 , solver='saga', penalty='l1' , tol=1e-3) # You can specify other paramete

# Fit the model on the training data
logreg.fit(X_train, y_train)

# Predict on the test set
predictions = logreg.predict(X_test)

# Evaluate the model
accuracy = logreg.score(X_test, y_test)
print(f"Accuracy: {accuracy:.2f}")

    Accuracy: 1.00
```



```

from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.datasets import load_iris # Using Iris dataset as an example

# Load the dataset (replace this with your dataset)
data = load_iris()
X = data.data # Features
y = data.target # Target

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the Support Vector Classifier
svm = SVC(kernel='rbf') # You can specify other parameters like kernel, C, gamma, etc.

# Fit the model on the training data
svm.fit(X_train, y_train)

# Predict on the test set
predictions = svm.predict(X_test)

# Evaluate the model
accuracy = svm.score(X_test, y_test)
accuracy_percentage = accuracy * 100
print(f"Accuracy: {accuracy_percentage:.2f}%")

    Accuracy: 100.00%

from sklearn.decomposition import PCA
from sklearn.svm import SVC

# Load the dataset (replace this with your dataset)
data = load_iris()
X = data.data[:, :2] # Considering only the first two features for visualization purposes
y = data.target

# Initialize the SVM classifier
svm = SVC(kernel='linear') # You can use different kernels

# Fit the SVM on the data
svm.fit(X, y)


# Use PCA to reduce dimensions for visualization (for plotting purposes)
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X)

# Create a meshgrid to plot decision boundary
x_min, x_max = X_pca[:, 0].min() - 1, X_pca[:, 0].max() + 1
y_min, y_max = X_pca[:, 1].min() - 1, X_pca[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.1),
                     np.arange(y_min, y_max, 0.1))

# Get predictions for meshgrid points
Z = svm.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

# Plot decision boundary
plt.contourf(xx, yy, Z, alpha=0.4)
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=y, edgecolor='k')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.title('SVM Decision Boundary Visualization (PCA)')
plt.show()

```



```

from sklearn.tree import DecisionTreeClassifier
# Using Iris dataset as an example

# Load the dataset (replace this with your dataset)
data = load_iris()
X = data.data # Features
y = data.target # Target

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the Decision Tree Classifier
clf = DecisionTreeClassifier(max_depth=3, random_state=42) # You can specify other parameters like max_depth, criterion

# Fit the model on the training data
clf.fit(X_train, y_train)

# Predict on the test set
predictions = clf.predict(X_test)

# Evaluate the model
accuracy = clf.score(X_test, y_test)
accuracy_percentage = accuracy * 100
print(f"Accuracy: {accuracy_percentage:.2f}")

Accuracy: 100.00

```

```

from sklearn.tree import DecisionTreeClassifier, plot_tree
import matplotlib.pyplot as plt

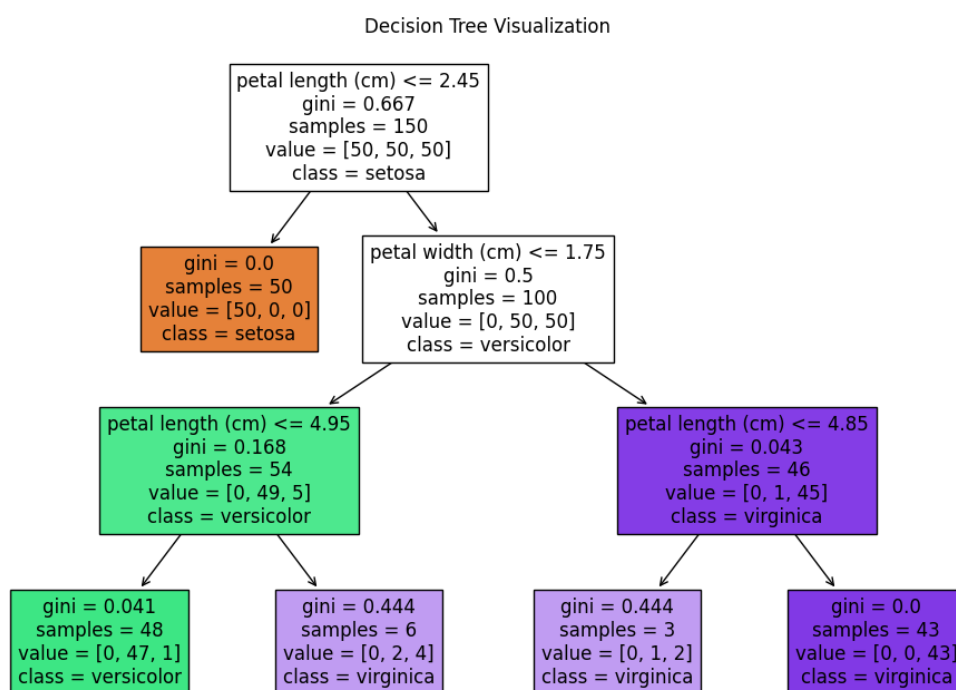
# Load the dataset (replace this with your dataset)
data = load_iris()
X = data.data # Features
y = data.target # Target

# Initialize the Decision Tree Classifier
clf = DecisionTreeClassifier(max_depth=3, random_state=42) # You can specify other parameters like max_depth, criter

# Fit the model on the entire dataset (for demonstration purposes)
clf.fit(X, y)

# Plot the Decision Tree
plt.figure(figsize=(12, 8))
plot_tree(clf, filled=True, feature_names=data.feature_names, class_names=data.target_names)
plt.title("Decision Tree Visualization")
plt.show()

```



```

from sklearn.datasets import fetch_california_housing
from sklearn.tree import DecisionTreeRegressor, plot_tree
import matplotlib.pyplot as plt

# Load the California housing dataset
data = fetch_california_housing()
X = data.data # Features
y = data.target # Target (continuous variable for regression)

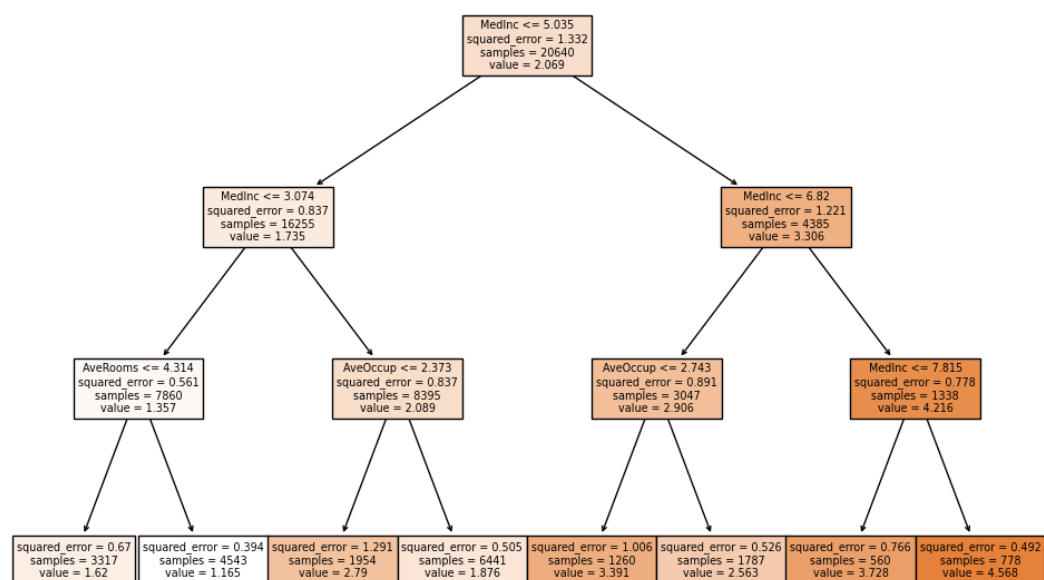
# Initialize the CART Regressor
regressor = DecisionTreeRegressor(max_depth=3, random_state=42) # You can specify other parameters like max_depth, c

# Fit the model on the entire dataset (for demonstration purposes)
regressor.fit(X, y)

# Plot the CART Decision Tree for regression
plt.figure(figsize=(12, 8))
plot_tree(regressor, filled=True, feature_names=data.feature_names)
plt.title("CART Decision Tree Visualization (Regression)")
plt.show()

```

CART Decision Tree Visualization (Regression)



```
from sklearn.datasets import make_classification
from sklearn.metrics import roc_curve, roc_auc_score

# Generate synthetic data for demonstration purposes
X, y = make_classification(n_samples=1000, n_features=20, n_classes=2, random_state=42)

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Initialize the Logistic Regression model
model = LogisticRegression()

# Fit the model on the training data
model.fit(X_train, y_train)

# Predict probabilities for the test set
y_prob = model.predict_proba(X_test)[:, 1]

# Calculate ROC curve and AUC score
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
auc = roc_auc_score(y_test, y_prob)

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f'AUC = {auc:.2f}')
plt.plot([0, 1], [0, 1], 'k--') # Diagonal line representing random guessing
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()
plt.show()
```



```
# Generate synthetic data for demonstration purposes
X, y = make_classification(n_samples=1000, n_features=20, n_classes=2, random_state=42)

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the Random Forest Classifier
rf = RandomForestClassifier(n_estimators=100, random_state=42) # You can specify other parameters like max_depth, et

# Fit the model on the training data
rf.fit(X_train, y_train)

# Predict on the test set
predictions = rf.predict(X_test)

# Evaluate accuracy
accuracy = accuracy_score(y_test, predictions)
accuracy_percentage = accuracy*100
print(f"Random Forest Accuracy: {accuracy_percentage:.2f}%")

Random Forest Accuracy: 90.00%
```

```
from sklearn.datasets import make_classification
from sklearn.ensemble import GradientBoostingClassifier

# Generate synthetic data for demonstration purposes
X, y = make_classification(n_samples=1000, n_features=20, n_classes=2, random_state=42)

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the Gradient Boosting Classifier
gb = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1, random_state=42) # You can specify other parame

# Fit the model on the training data
gb.fit(X_train, y_train)

# Predict on the test set
predictions = gb.predict(X_test)

# Evaluate accuracy
accuracy = accuracy_score(y_test, predictions)
accuracy_percentage = accuracy*100;
print(f"Gradient Boosting Accuracy: {accuracy_percentage:.2f}%")

Gradient Boosting Accuracy: 91.00%
```

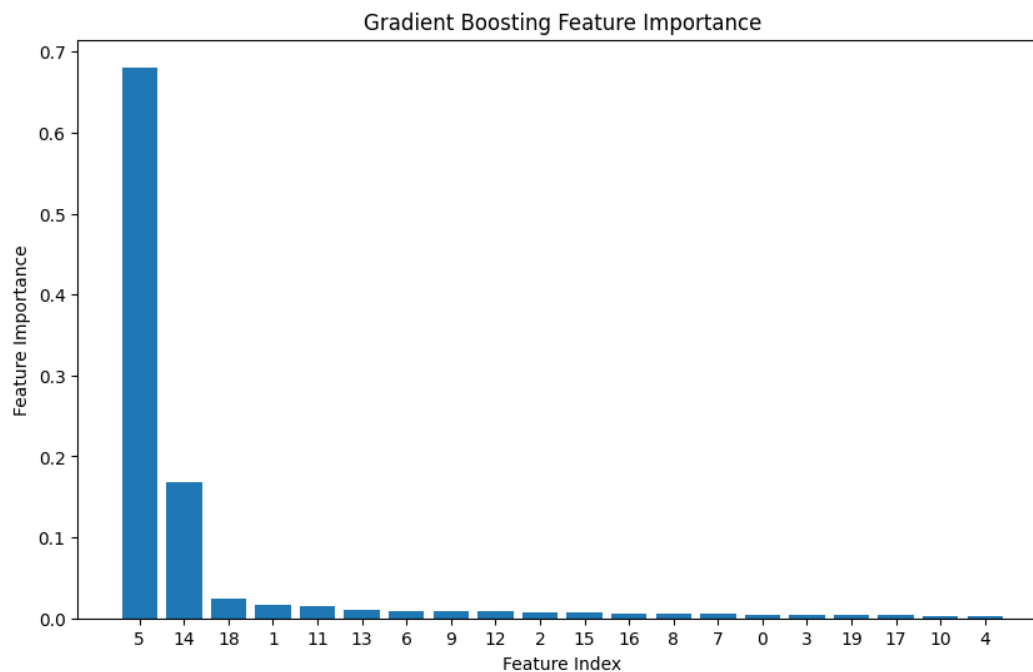
```
# Generate synthetic data for demonstration purposes
X, y = make_classification(n_samples=1000, n_features=20, n_classes=2, random_state=42)

# Initialize the Gradient Boosting Classifier
gb = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1, random_state=42)
gb.fit(X, y)

# Get feature importances
feature_importance = gb.feature_importances_

# Sort feature importances in descending order
indices = feature_importance.argsort()[::-1]

# Plot feature importances
plt.figure(figsize=(10, 6))
plt.title("Gradient Boosting Feature Importance")
plt.bar(range(X.shape[1]), feature_importance[indices], align="center")
plt.xticks(range(X.shape[1]), indices)
plt.xlabel("Feature Index")
plt.ylabel("Feature Importance")
plt.show()
```



```

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.datasets import load_iris # Using Iris dataset as an example

# Load the dataset (replace this with your dataset)
data = load_iris()
X = data.data # Features
y = data.target # Target

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the Random Forest Classifier
rf = RandomForestClassifier(n_estimators=100, random_state=42) # You can specify other parameters like n_estimators,

# Fit the model on the training data
rf.fit(X_train, y_train)

# Predict on the test set
predictions = rf.predict(X_test)

# Evaluate the model
accuracy = rf.score(X_test, y_test)
accuracy_percentage = accuracy * 100
print(f"Accuracy: {accuracy_percentage:.2f}%")

    Accuracy: 100.00%

# Create NumPy arrays from the columns
recency = np.array(data['Recency'])
kidhome = np.array(data['Kidhome'])

# Compute Pearson correlation coefficient using NumPy
correlation_matrix = np.corrcoef(recency, kidhome)

# Extract the correlation coefficient
correlation = correlation_matrix[0, 1]

print(f"Pearson Correlation for recency and kidhome: {correlation}")

    Pearson Correlation for recency and kidhome: -1.2819751242557092e-17

print(data.columns)

    Index(['ID', 'Year_Birth', 'Education', 'Marital_Status', 'Income', 'Kidhome',
        'Teenhome', 'Recency', 'MntWines', 'MntFruits', 'MntMeatProducts',
        'MntFishProducts', 'MntSweetProducts', 'MntGoldProds',
        'NumDealsPurchases', 'NumWebPurchases', 'NumCatalogPurchases',
        'NumStorePurchases', 'NumWebVisitsMonth', 'AcceptedCmp3',
        'AcceptedCmp4', 'AcceptedCmp5', 'AcceptedCmp1', 'AcceptedCmp2',
        'Complain', 'Z_CostContact', 'Z_Revenue', 'Response'],
        dtype='object')

# Basic filter method: Filtering rows based on a condition
# For example, filtering rows where 'Column_A' values are greater than 50
filtered_df = data[data['Recency'] > 98]

# Display the filtered DataFrame
print("\nFiltered DataFrame:")
print(filtered_df)

    Filtered DataFrame:

```


	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	\
38	8595	1973	Graduation	Widow	42429.0	0	
192	7829	1900	2n Cycle	Divorced	36640.0	1	
208	868	1966	Graduation	Married	44794.0	0	
444	2106	1974	2n Cycle	Married	20130.0	0	
491	22	1976	Graduation	Divorced	46310.0	1	
606	7232	1973	Graduation	Widow	42429.0	0	
685	10142	1976	PhD	Divorced	66476.0	0	
700	9977	1973	Graduation	Divorced	78901.0	0	
725	7212	1966	Graduation	Married	44794.0	0	
1033	5263	1977	2n Cycle	Married	31056.0	1	
1171	3363	1974	2n Cycle	Married	20130.0	0	
1473	4070	1969	PhD	Married	94871.0	0	
1542	528	1978	Graduation	Married	65819.0	0	
1685	7947	1969	Graduation	Married	42231.0	1	
1800	2831	1976	Graduation	Together	78416.0	0	
1820	2415	1962	Graduation	Together	62568.0	0	
1894	1743	1974	Graduation	Single	69719.0	0	

	Teenhome	Recency	MntWines	MntFruits	...	NumWebVisitsMonth	\
38	1	99	55	0	...	5	
192	0	99	15	6	...	5	
208	1	99	54	0	...	6	
444	0	99	0	6	...	8	
491	0	99	185	2	...	8	
606	1	99	55	0	...	5	
685	1	99	372	18	...	4	
700	1	99	321	11	...	4	
725	1	99	54	0	...	6	
1033	0	99	5	10	...	8	
1171	0	99	0	6	...	8	
1473	2	99	169	24	...	7	
1542	0	99	267	38	...	3	
1685	1	99	24	0	...	5	
1800	1	99	453	38	...	3	
1820	1	99	362	17	...	4	
1894	0	99	273	86	...	1	

	AcceptedCmp3	AcceptedCmp4	AcceptedCmp5	AcceptedCmp1	AcceptedCmp2	\
38	0	0	0	0	0	
192	0	0	0	0	0	
208	0	0	0	0	0	
444	0	0	0	0	0	
491	0	0	0	0	0	
606	0	0	0	0	0	
685	0	0	0	0	0	
700	0	0	0	0	0	
725	0	0	0	0	0	
1033	0	0	0	0	0	
1171	0	0	0	0	0	
1473	0	1	1	0	0	
1542	0	0	0	0	0	
1685	0	0	0	0	0	
1800	0	0	0	0	0	
1820	0	0	0	1	0	
1894	0	0	0	0	0	

```

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

```

```
# Assuming 'X' contains your features and 'y' contains the target variable
```

Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.